

My Titled

Andreas Lillevang Bech

June 2019

Abstract

This document provides an introduction to R Markdown, argues for its...

1 Introduction

This paper revolves around replicating Gu, Kelly, and Xiu (2018) in which return the predictability on financial assets is explored. A catalog of Machine Learning methods are evaluated on their relative performance on varying time horizons and data generating processes, which is done by a simulation study.

The paper is structured as follows. A theoretical section introduces the statistical models under evaluation as well as the sample splitting and choosing between model specifications. Then the results of the simulation study are presented. Finally the methods are tested on a small empirical example.

2 Theoretical section

2.1 Statistical Model

The model most often considered in the context of statistical learning is the additive error model.

$$r_{i,t+1} = E_t(r_{i,t+1}) + \epsilon_{t+1}$$

where

$$E_t(r_{i,t+1}) = g(z_{i,t})$$

The conditional expectation, which is the aim of the estimation procedure, is a function only of the available predictors at time t , $z_{i,t}$. In a statistical setting a pair $(r_{i,t+1}, z_{i,t})$ will not have a deterministic relationship of the form $r_{i,t+1} = E_t(r_{i,t+1})$. Therefore an additive error, ϵ_{t+1} , is usually added to give a complete characterization of the truth¹.

The difference in this setting, as compared to the usual cross-sectional machine learning framework, is the time component. The data in this setting is a panel of stocks, each stock indexed by i and the time period by t . Thus predictors at time t of stock i , $z_{i,t}$, may include past observations of the response or predictor variables, e.g recent price trends like reversal and momentum². This adds some considerations for the estimation procedure. The conditional expectation is a function of the predictors $z_{i,t}$ and here it is denoted g . g is assumed to have the same form across time the individual stock which gives in all $N \times T$ observations for estimation, where T is the no. of time periods and N is the number of individual stocks. However, there is a time order to the observations which must be preserved due to the nature of the predictors. This restricts the sample-splitting procedure slightly.

¹See Friedman, Hastie, and Tibshirani (2001)

²Which may or may not be proxies for time-varying beta compensation Kelly, Moskowitz, and Pruitt (2018)

2.2 Tuning parameters (hyper parameters)

Machine learning models have so-called tuning parameters which usually qualify the degree of complexity of the models, i.e. how much they try to adapt to a given set of training data. Examples of tuning parameters are the penalty parameter in Lasso estimation, the no. of knots and degrees in a piece wise polynomial function, or the depth of a regression tree. As the training error always decreases with more complexity, a “validation” procedure is needed to restrict the model. Validation amounts to estimating the test error in order to choose the degree of complexity in a given model that performs best for forecasting/prediction.

Due to the time element in the data, which we wish to conserve in chronological order (some predictor variables like momentum use the past returns), cross-validation is not appropriate for choosing tuning parameters. Instead, sample-splitting, dividing the training data into a *training* and a *validation* part will allow to evaluate an estimated model on new data while still keeping the test data locked up for final evaluation. Choosing tuning parameters simply amounts to choosing the model that performs best in the validation sample. The following pseudo code illustrates the idea.

```
# split data
data = ...
train, validation, test = split(data)

# tune model hyperparameters
parameters = ...
for params in parameters:
    model = fit(train, params)
    skill = evaluate(model, validation)

# evaluate final model for comparison with other models
model = fit(train)
skill = evaluate(model, test)
```

2.3 Linear models

Linear regression is a model based approach which imposes that the conditional expectation is linear in its arguments $f(x) = x'\theta$.

For the **simple linear model** where each predictor enters as an argument as it is this gives $g(z_{i,t}) = z'_{i,t}\theta$. The parameters, θ of the model can be estimated using the training data, most popularly by *least squares*, which picks the parameters to minimize the objective function of residual sum of squares

$$RSS(\theta) = \sum_{i=1}^N \sum_{t=1}^T (r_{i,t+1} - g(z_{i,t}; \theta))^2$$

Here there are $T \times N$ observations as the pooled estimate of θ over the panel of individual stocks is considered.

It is well known that the linear model has a low bias. However, in the face of a large amount of predictors relative to number of observations, the linear model will have large variance. In this context, as Gu, Kelly, and Xiu (2018) points out, what really matters is the no. of predictors p vs. time periods T , as stock returns typically share a large cross-sectional dependence which limits the information gained by adding observations along the cross-section.

2.3.1 Penalized Regression

Although there are variable subset select procedures that can be performed even on large models, like Forward- and Backward-Stepwise selection³, a more efficient way to reduce the complexity of the linear model is to impose constraint on the size of θ . This is done via a penalty in the objective function, such that parameters are estimated by minimizing

$$\mathcal{L}(\theta, \cdot) = RSS(\theta) + \phi(\theta, \cdot)$$

where the penalty function ϕ depends on a tuning parameter.

For **Ridge regression** an L2 penalty is applied to the parameters. In this case $\phi(\theta, \lambda) = \lambda \|\theta\|_2^2$ such that the parameters are constrained in a hypersphere in the parameter space with center at the origin. Shrinking the parameters in this way introduces bias to the linear model but can be shown to reduce the variance. In terms of expected prediction error this often leads to better out of sample predictions in high dimensional problems.

Another proposed way of shrinking the parameters is with an L1 penalty, $\phi(\theta, \lambda) = \lambda \|\theta\|_1$. Regression with this penalty is known as **Lasso regression**. Geometrically this means that the parameters are constrained in a “diamond” (L1 unit circle), which will tend to set many parameters to exactly zero and thus working as actual subset selection of the variables. This is opposed to the Ridge regression that will never set parameters to zero.⁴

Finally a combination of the two approaches to the penalty has been proposed by Zou and Hastie (2005), known as the **Elastic Net**. It uses a linear combination between the penalties, $\phi(\theta, \lambda, \alpha) = \lambda (\alpha \|\theta\|_2^2 + (1 - \alpha) \|\theta\|_1)$. For $\alpha < 1$ the constraint boundary still has “kinks” at $\theta_j = 0$ and thus still enjoys the property of the Lasso of deselecting some variables.

2.3.2 PCR and PLSR

Principal Components Regression and Partial Least Squares Regression are linear models that consider certain linear combinations of the predictors and use these as arguments in the linear model. Such methods are derived to reduce the dimension of the problem, which is especially useful if there are a large number of highly correlated predictors. The two methods differ in how they utilise the inputs.

PCR seeks to use as much variation from the inputs in as few variables as possible. It does this by using the principal components of the input data. The largest principal component is the direction in the input space that yields the largest variance if all input point where projected onto this line. Call a vector point in this direction $v_1 \in \mathbb{R}^I$, then the first/largest principal component is Xv_1 . The second principal component maximises this variance subject to being orthogonal to v_1 . There are in total P principal components and the choice of how many to include in the linear model as arguments is tuning parameter subject to validation.

The second approach, **PLSR**, instead tries to maximise the covariance with the responses (returns). In essence the first argument, the first partial squares direction, put forward to the model is a weighted sum of the predictors where predictors (after having been standardized) with higher covariance with the response have higher weights. To include more directions the inputs are orthogonalized (by gram-schmidt) and the procedure is repeated again. How many arguments to give the model is again a tuning parameter.

2.3.3 Generalized Linear Model

The final linear model considered in this exercise is the **GLM**. Each predictor is transformed by a *linear basis expansion* and a linear model is then fit to the new space of derived inputs. Because the basis functions most

³see Friedman, Hastie, and Tibshirani (2001), chapter 3.3

⁴see again Friedman, Hastie, and Tibshirani (2001), chapter 3.3, for a visual illustration of the difference between Ridge and Lasso penalty.

often are non-linear such a model can capture non-linearities in the data with the regression function/surface $f(x)$ being non-linear. However the model is still linear in the coefficients to be estimated which is such a model is also referred to as “additive”.

In this exercise a second order spline expansion is chosen, as in Gu, Kelly, and Xiu (2018). A quadratic spline is a piecewise/local second order polynomial but where it is restricted to be continuous and have a continuous first derivatives at the “knots”. Having restricted the local polynomials to be quadratic, the only tuning parameter is the no. of knots, more knots leading to a more complex model.

As the number of parameters in such a model is much greater than the simple linear model it is surely advantageous to use a shrinking method. In this setting the Group Lasso is useful. It gives the penalty to the objective function

$$\phi(\theta, \lambda, K) = \lambda \sum_{j=1}^P \left(\sum_{k=1}^K \theta_{j,k}^2 \right)^{1/2}$$

where P is the no. of predictors available and K is the no. of terms in a given spline considered. The Group Lasso penalty has the desirable property that terms based on the same predictor are set to zero together. Thus for the GLM considered here, the no. of knots and the penalty parameter λ are the tuning parameters.

2.4 Tree-based Models

Tree regression takes completely different approach to approximating the conditional expectation of the response given the inputs, $f(x) = E[Y|X = x]$. Instead of using a continuous surface over the predictor space the approach is instead to divide the predictor space into rectangles and treat $f(x)$ as constant in each rectangle. This is a very direct approach to approximate $E[Y|X = x]$, where the (empirical) mean is conditions not at the point x but in on an area around x . The size of each rectangle is what determines the complexity of a tree regression.

To find the optimal division of the predictor space by considering all options is computationally infeasible. The most common procedure is a *greedy* one where at each step the space or a “branch” of the space is split by a predictor and a split value, both of which are chosen to reduce the RSS maximally for that step. Each split creates two new branches. This procedure known as *recursive binary splitting*.

Splitting/branching halts when the tree reaches a prespecified stopping criterion. In this paper the stopping criterion is based on a minimal no. of observations in each terminal node (“leaf”) or reach of a depth limit, both are chosen as tuning parameters based on the validation sample.

Unlike the GLM where individual predictors can predict in a non-linear fashion, regression trees allow for a more flexible non-linear effects, particularly in the form interactions between variables. A tree divides the p -dimensional input space into (hyper)rectangles which allows for interactions for up to all p predictors. A regression tree, however, suffers from high variance in the sense that a small perturbation in the input values can lead to a large change in a given tree. This amounts to easily overfitting on the training data. The variance issue of trees has been very successfully addressed by forecasting based on averaging over many trees, a method known as ensemble learning, where typically simple/shallow trees are used. Two ways of this are considered approach, “boosting” and “bagging”.

Boosting fits an oversimplified/“shallow” tree on the training data and then fits a consecutive tree on the residuals from the first tree. Then the sum of the forecasts for the two trees is considered as the forecast, but where the second tree’s forecast (of the residuals) is shrunk by a factor between $\nu \in (0, 1)$. A third tree is added to these residuals and so on.. Tuning parameters to be chosen are the limit depth of each tree, the shrinkage factor and the total number of trees to fit. For boosting a maximal of 6 splits is typically considered, hence the “slow learning” approach.

For the bagging, many trees are fitted using (nonparametric-)bootstrap samples and then averaged. To avoid using the most influential variables for the first splits in all the trees, only a random subset of the

predictors are considered for splitting at each branch. This amounts to a **Random Forest**. The Random Forest is fairly straightforward to use “out-of-the-box” and performs very well without having a lot of tuning parameters to be chosen. The most important tuning parameter is the fraction of predictors available for each split. The number of trees should be chosen but the MSE will stabilize after a while. The node size should also be chosen, as a stopping criterion, which also represents a bias-variance trade-off as smaller nodes will lead to more complex trees.

2.5 Neural Networks

The final learning method considered, which is in a category of its own, is the **Neural Network** (NN). The inputs are filtered through “hidden layers” to create new arguments to use as predictors in a linear model. This is analogous to PCR and PLSR, however, the neural network is not as systematic and is much more flexible. Figure 1 gives a simple illustration of a Neural Network in the regression context. There are 4 input/predictors and one hidden layer with 3 “neurons”. Each neuron gets its value as linear combination of all the neurons from the previous layer. Without any hidden layer this is just a simple linear regression.

The NN allows for non-linear regression by applying a non-linear “activation function” to each neuron in the hidden layers. In fact a NN with just a single hidden layer containing a finite number of neurons can approximate continuous functions on compact subsets of \mathbb{R}^2 (see Cybenko 1989). The NN’s in this exercise uses the ReLU activation function $ReLU(x) = \max\{x, 0\}$. The NN’s flexibility is also one of its greatest criticisms and it is often described as a “black box” without any meaning, just pure pattern recognition.


For the number of parameters in a neural network, consider a NN with one hidden layer, with, say, n neurons. There are p predictors, thus each neuron in the hidden layer has to it associated $(p+1)$ parameters (it’s a linear model in itself). This gives $n \times (p+1)$ parameters in the hidden layer. Additionally, there will be $(n+1)$ parameters in the output layer which constitute the linear model for predicting the output with the neurons from the hidden layer. In all $n \times (p+1) + (n+1)$ parameters in the neural network with one hidden layer containing n neurons. The large number of parameters are obviously a problem for contexts with few observations for estimation, like in the context of economics and finance. Thus in these contexts shallow networks with few neurons in each tend to perform better.

In this exercise NN’s with layers ranging from 1-3 are considered, and the no. of neurons in each layer is set by the geometric pyramid rule (see Masters 1993). In estimation RRS is used as an objective. It is optimized by Stochastic Gradient Descent (SGD) and an l1 penalty on the parameters is used to prevent overfitting - this is also called “weight decay”. Another measure against overfitting is “early stopping” which evaluates the error on a validation set after each epoch in the SGD. If the validation error increases for 3 epochs in a row the fitting is halted. Important tuning parameters are the parameter penalties and the step size in the SGD.

2.6 Evaluation

In this paper model performance is based on the R-squared measure. Both for in-sample and out-of-sample the RRS is compared to the total sum of squares based on the historical mean of the returns, i.e the average of returns before the test sample. Out-of-sample R-squared is then

Figure 1: Illustration of simple Feed-forward Neural Network



$$R_{OOS}^2 = 1 - \frac{\sum_{(i,t) \in \text{test}} (r_{i,t+1} - \hat{r}_{i,t+1})^2}{\sum_{(i,t) \in \text{test}} (r_{i,t+1} - \bar{r}_{train})^2}$$

Models can be compared pairwise for their out-of-sample predictive ability by the Diebold-Mariano test⁵. The test considers the difference in errors

$$d_t = \left(e_{i,t}^{(1)}\right)^2 - \left(e_{i,t}^{(2)}\right)^2$$

where errors $e_{i,t}$ are transformed with the square - the absolute value could also be used⁶. Under the null hypothesis both model perform equally well, $\bar{d} = 0$ and the limiting distribution is

$$\sqrt{T}\bar{d} \xrightarrow{d} N\left(0, \sum_{k=-\infty}^{\infty} \gamma_d(k)\right)$$

given that this second moment is bounded. The long run variance of the sample mean is⁷

$$\text{Var}(\sqrt{T}\bar{d}) \xrightarrow{T \rightarrow \infty} \sum_{k=-\infty}^{\infty} \gamma_d(k)$$

Thus test statistics will be

$$\frac{\bar{d}}{\sqrt{\frac{\sum_{k=-\infty}^{\infty} \gamma_d(k)}{T}}}$$

which is asymptotically standard normal. The HAC estimator proposed by Newey and West (1987) is used to get a consistent estimate of the variance.

3 Simulation

3.1 Data Generating Process

The simulation is based on the experiment outlined in Gu, Kelly, and Xiu (2018). The experiment is set up to compare model performance for different data generating processes.⁸

DGP-a is a linear DGP such that returns in the following period for individual stock i is a linear combination of three of the respectively 100 and 200 available predictors. To the extent that machine learning models is thought to capture non-linearities among the predictors, the simple linear model (OLS)⁹ should perform relatively well in this. A different number of total predictors for each stock is considered to evaluate the improvement shrinking and dimensionality reductions methods can provide.

⁵Diebold and Mariano (2002)

⁶If doing the test on a panel it makes to average prediction error for each model in each time period, i.e. in the cross-section. This is due to the high error dependence across stocks, which likely violates the conditions for asymptotic normality of the test statistic.

⁷see Hamilton (1995)

⁸I will give the outline of the experiment. The specifics can be found in the appendix of Gu, Kelly, and Xiu (2018) under the name "Monte Carlo simulations".

⁹Here and in the rest of the paper OLS is meant to be the naive linear model of all predictors entering without transformations. No variable selection is performed. Other linear models considered in the simulation study are the penalty methods of Ridge, Lasso and Elastic net, and the Generalized Linear Model (Generalized Additive Model) where splines of variables are used. The oracle is the linear model where only the true predictors from the DGP enter

The other DGP considered, DGP-b is a non-linear counterpart. The same predictors enter the true DGP in the construction of returns/responses, however, each variable enter as a non linear transformation.

Specifically the setup is like this. The returns are made up by the conditional expectation, g , given the predictors, plus an error.

$$r_{i,t+1} = g(z_{i,t}) + e_{i,t+1}$$

$e_{i,t+1}$ is a combination of a normal error and a student error, it is common to both DGP's¹⁰

Choosing the data generating process is choosing g . The set of predictors for each stock i , $z_{i,t}$ is made up of individual characteristics, $c_{i,t}$ and interactions of the characteristics with a common macroeconomic variable, x_t . That is, $z_{i,t} = (1, x_t)' \otimes c_{i,t}$. If there are Pc individual characteristics there will thus be $P = 2 \times Pc$ predictors for each individual stock i at time t . The individual characteristics are generated using an AR(1) process with normal error and with AR-coefficient randomly distributed uniform(0.9, 1), such that there is a high degree of persistence in the characteristics. The macroeconomic variable, x_t is similarly generated as an AR(1) with normal error and AR-coefficient=0.95.

$$\begin{aligned} (a) &= g(z_{i,t}) = (c_{i1,t}, c_{i2,t}, c_{i3,t} \times x_t) \theta_0, \quad \text{where } \theta_0 = (0.2, 0.2, 0.2)' \\ (b) &= g(z_{i,t}) = (c_{i1,t}^2, c_{i1,t} \times c_{i2,t}, \text{sgn}(c_{i3,t} \times x_t)) \theta_0, \quad \text{where } \theta_0 = (0.04, 0.03, 0.012)' \end{aligned}$$

3.2 Results

Due to a computing power issue a slight deviation from the simulation study in Gu, Kelly, and Xiu (2018) is that the number of firms (N) is 100 instead of 200 and the number of periods (T) is 90 instead of 180. It is not clear in the paper whether or not an expanding window with recursive estimation is considered in the simulation study. Due to the computing power issue a fixed window is used here where only one estimation is needed for each simulation repetition. This decision should not have a large impact as the DGP remains the same over “time” in the simulated data.

The linear models to be compared are pooled OLS, Ridge, Lasso and Elastic Net (ENET). OLS is the naive approach where variable selection is not performed. Further the more complex linear models of Principal Component regression (PCR), Partial Least Squares regression (PLSR) and the Generalized Linear/Additive model (GLM) are also estimated on each simulated data. For tree based methods, Gradient Boosted regression trees (GBRT) and Random Forest are considered. Finally Neural Networks (NN) with 1-3 layers are estimated.

The results of the simulation study can be found in Table 1. First of all, the Neural Networks were not trained correctly as some did not converge, at all, to a minimum¹¹, therefore these repetitions were removed when calculating the MC-average.

Second, the training of the OLS model and the Oracle might have been unfair. As there are no tuning parameters in these models they were instead estimated on the whole training sample, training plus validation, and then used for prediction on the test sample. Given that it is easier to predict longer horizons with this kind of data, this decision might have put these models at a disadvantage as compared to the other models, which were only estimated on the training sample, without the validation sample.

In terms of replicating Gu, Kelly, and Xiu (2018) it is mostly a positive result. For the linear DGP, (a), the Ridge and Elastic net regressions perform very well. This is unsurprising as these are linear models with shrinking and variable selection, allowing them to come potentially close in model specification to the true process. The dimension reduction methods, like PCR and PLSR, do not perform as well. Likely due to the very low signal-to-noise ratio in this kind of data. The very complex and adaptive linear model, GLM, does

¹⁰see Gu, Kelly, and Xiu (2018) for details.

¹¹The NN has to option of, for example, setting all weights in the hidden layer to zero, which would result in just a reasonable loss. It is suspected the the problem of the design of the optimization algorithm lies in the “patience” parameter in the “early-stopping” mechanism. The patience was set to 2 while this might be too little considering

Table 1: Comparison of Predictive R^2 s for Machine Learning Algorithms in Simulations

DGP	(a)				(b)			
Parameter	$P_c = 50$		$P_c = 100$		$P_c = 50$		$P_c = 100$	
$R^2(\%)$	IS	OOS	IS	OOS	IS	OOS	IS	OOS
OLS	10.11	-1.30	12.90	-2.76	6.87	-5.14	9.45	-7.64
Ridge	8.06	1.47	9.35	1.01	4.56	-0.19	5.54	-0.15
Lasso	8.12	2.61	9.37	1.86	4.54	-0.43	5.60	-0.03
ENet	8.18	2.69	9.53	1.89	4.66	-0.41	5.72	-0.07
PCR	7.46	0.08	8.32	-0.04	4.23	-0.68	4.97	-0.72
PLSR	7.38	0.27	8.51	-0.16	3.67	-1.98	4.50	-2.37
GLM	7.92	0.21	9.07	0.01	4.80	-0.48	5.87	-0.18
RandomF	13.91	2.33	16.39	1.83	13.49	1.02	14.74	2.23
GBRT	9.84	1.94	11.59	1.32	10.06	1.09	10.37	1.20
NN1	1.65	0.74	-0.28	0.81	1.41	-0.07	-2.96	0.00
NN2	7.15	0.15	2.26	0.26	13.08	-0.10	-3.57	-0.07
NN3	6.17	-8.25	0.18	-2.65	3.35	-0.18	4.56	-1.91
Oracle	8.41	3.45	10.04	3.74	4.89	0.12	5.81	-0.07

Note: The table report in-sample (IS) and out-of-sample (OOS) R^2 s for the different data generating processes (a) and (b). The forecasting methods use are OLS using all available variables, Ridge, Lasso, Elastic Net (ENet), Principal Component Analysis (PCR), Partial Least Squares (PLSR), Random Forest, Gradient Boosted Tree Regression (GBRT), and Neural Networks with 1-3 layers (NN). The Oracle is a linear model using only variables present in the true DGP. $N=100$, $T=90$ and P_c considered are 50 and 100. The number of Monte Carlo repetitions is 50.

not perform well either, and probably for the same reason. The GLM has a lot of variables to estimate in a linear regression and thus suffers from high variance. The tree based regressions perform as well as the best linear ones, even though the DGP is completely linear. This goes in particular for the Random Forest.

The more interesting finding is to be made for the non-linear DGP, (b). Recall that in this case the returns are constructed by the square of an input, an interaction between two inputs and the sign function of an input, which is of course not continuous. Now almost all models perform worse than the historical average by having a negative R^2 out-of-sample. The only models to perform with a positive R^2 are the tree based models. This is in line with Gu, Kelly, and Xiu (2018) and points towards these models, *homogeneous* ensembles of trees, as being very effective in high noise environments. In fact, GBRT and Random Forests perform even better with $P_c = 100$ than for $P_c = 50$, i.e. with more noise, both here and in the paper by Gu, Kelly, and Xiu (2018).

Table 2 explores the relative performance of the different methods over longer horizons. For DGP (a), as compared to 1-period out-of-sample, performance is *better* for models that performed well in 1-period forecasts, but *worse* for models that performed poorly before. As an example, the shrinkage and selection methods of Lasso and Elastic Net, which performed well for the linear DGP, (a), perform even better on longer horizons. For the annual horizon these models explain almost 15% of the out-of-sample variation. Likewise, the tree based regressions perform better the longer the horizon with the Random Forest being able to explain 20% of the variation of the linear process one year out of sample. In contrast, the poorer models perform similarly or worse on longer horizons. The Neural Network performs extremely bad on long horizons meaning that it fails to pick up on the essential part of the DGP, i.e. it finds a “valley” far away from the true minimum.

For the non-linear DGP, (b), all models perform worse, except for the Random Forest. It seems that longer horizons does not aid in finding the true DGP despite the underlying variables, the characteristics, being generated as mean-reverting processes (AR(1)’s). This does not go, however, for the Random Forest, which performs better the longer the horizon. An explanation for this is that the trees used in Random Forest are complex and thus has the ability to pick up on a non-linear, discontinuous pattern, while still avoiding

Table 2: Comparison of Predictive R^2 s for Machine Learning Algorithms in Simulations

DGP	(a)						(b)					
Horizon	Quarter		Halfyear		Annual		Quarter		Halfyear		Annual	
$R^2(\%)$	IS	OOS	IS	OOS	IS	OOS	IS	OOS	IS	OOS	IS	OOS
OLS	22.45	-5.49	29.91	-16.92	44.61	-10.61	15.30	-20.43	25.37	-28.22	33.69	-44.38
Ridge	14.39	2.81	18.02	1.08	29.24	5.54	6.01	-0.63	12.00	-0.48	15.44	-4.27
Lasso	14.27	6.00	17.20	7.88	27.92	14.84	5.20	0.09	11.09	-0.28	11.81	-2.86
ENet	14.52	6.11	17.76	8.03	28.40	14.86	6.09	-0.42	11.87	-0.10	14.65	-4.26
PCR	11.98	0.23	15.31	-5.32	22.99	0.94	3.94	-1.13	9.84	-1.30	11.61	-5.16
PLSR	11.63	0.94	12.05	-4.45	27.24	-1.61	6.41	-7.93	9.72	-8.70	13.72	-20.01
GLM	13.46	0.30	17.54	-4.13	29.03	0.12	6.76	-2.44	10.12	-2.43	11.77	-6.03
RandomF	88.03	6.64	91.13	10.27	94.29	20.30	87.51	4.70	91.34	7.58	93.72	10.04
GBRT	17.84	5.99	21.89	7.67	36.57	12.07	15.67	2.51	22.53	2.90	30.81	-1.57
NN1	6.68	5.22	-0.38	-21.73	23.42	-61.45	0.23	-4.55	-3.66	-34.08	9.07	-71.39
Oracle	15.33	8.91	18.72	10.93	28.75	20.59	6.30	-0.48	11.48	-0.48	11.46	-1.25

Note: The table reports in-sample (IS) and out-of-sample (OOS) R^2 s for the different data generating processes (a) and (b) and longer horizons. That is 3 months, 6 months and 12 months. The data frequency is monthly. The forecasting methods use are OLS using all available variables, Ridge, Lasso, Elastic Net (ENet), Principal Component Analysis (PCR), Partial Least Squares (PLSR), Random Forest, Gradient Boosted Tree Regression (GBRT), and Neural Networks with 1-3 layers (NN). The Oracle is a linear model using only variables present in the true DGP. N=100, T=90 and P=100. The number of Monte Carlo repetitions is 50.

overfitting with the use of bagging.

4 Empirical Exercise

Given data compiled for the paper Welch and Goyal (2008)¹². The sample considered runs from 1947-03-01 to 2017-06-01. The data is monthly in frequency but some variables are only updated on a quarterly or annual basis. Index is the S&P 500 w/o dividends.

The data is divided into training, validation and test samples. The training sample runs in 1947-1982, 35 years, the validation sample in (1982-2000), 18 years, and finally the test sample the final 17 years.

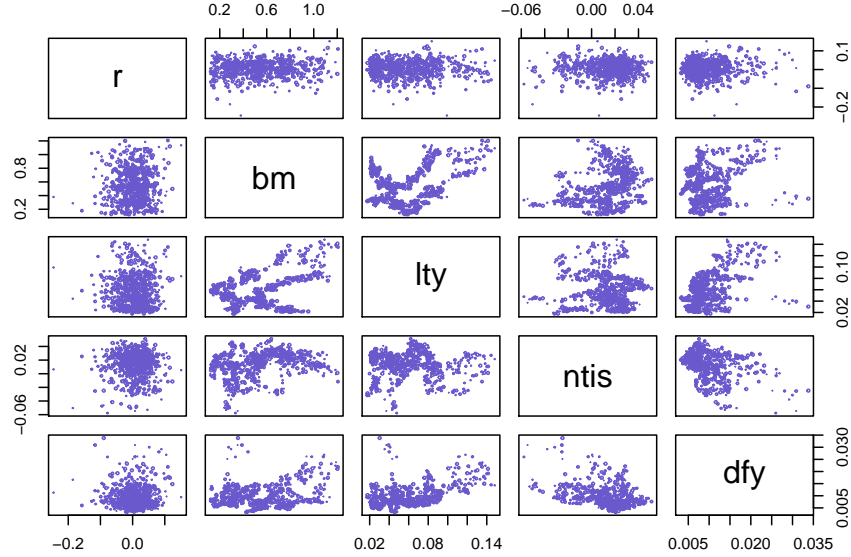
The crucial difference between the data used in the paper Gu, Kelly, and Xiu (2018, 2018) and the data collected here is the number of observations. In the data from Welch and Goyal (2008) only returns on the stock market index (S&P500) is considered as the response. Thus, the only relevant predictors are aggregated macroeconomic characteristics from the stock market including dividend-price ratio (dp), earnings-price ratio (ep), book-to-market ratio (bm), net equity expansion (ntis), Treasury-bill rate (tbl), term spread (tms), default spread (dfy), and stock variance (svar). In Gu, Kelly, and Xiu (2018) data are obtained at an individual stock level and with stock individual characteristics, such as, growth in common shareholder equity and gross profitability, can therefore be used as predictors. This gives more observation as well as more predictors, which in turn gives a more stable estimation and more information to forecast with. Nevertheless it is an interesting exercise to see how the different machine learning models compare in the case considered here (with no cross-section).

Note: **r** is log returns on the S&P500, **bm** is the ratio of book value to market value for the Dow Jones Industrial Average, **lty** are the Long-term government bond yields from Ibbotson's Stocks, Bonds, Bills and Inflation Yearbook¹³, **ntis** is the ratio of twelve-month moving sums of net issues by NYSE listed stocks

¹²The data, including a appendix with sources and description of the data, can be found on Amit Goyal's website <http://www.hec.unil.ch/agoyal/>

¹³See Stocks (1995)

Figure 2: Scatterplot-matrix for the data from @goyalwelch2008



divided by the total market capitalization of NYSE stocks and finally **dfy** the default yield spread is the difference between BAA- and AAA- rated corporate bond yields.

A scatterplot matrix of the data, Figure 2, shows well the small correlation between the log returns and the predictors. In Figure 3 are included the log returns as well as the 4 most relevant predictors in terms of significance in correlation. Looking at the top row it is clear that there is very little “signal” to extract from the data, at least from linear dependence.

Table 3: Comparison of Predictive R^2 s for Machine Learning Algorithms in Simulations

	OLS	Ridge	Lasso	ENet	PCR	PLSR	GLM	RandomF	GBRT
OOS	-21.60	1.22	-2.51	-3.16	-7.21	-13.55	-19.95	-9.59	-3.60
IS	9.29	3.70	3.69	3.70	4.70	-1.08	3.69	12.42	9.89

Note: Model performance on the data from @goyalwelch2008. Neural Networks are excluded as they did not come to a reasonable estimate. Data can be found at Amit Goyal’s website <http://www.hec.unil.ch/agoyal/>.

Table 3 gives the performance of the models on the empirical data. Out-of-sample all models perform worse than the historical average, except the Ridge regression which performs slightly better. Comparing the in-sample forecasts the Ridge has among the highest error, which indicates that the Ridge does well not to overfit on the data. This property is useful in an environment with not many observations to train on and very little signal in each observation. The tree based regressions which performed very well in the simulations perform significantly worse now, with the Random Forest being among the poorest out-of-sample predictors.

Model performance can be tested pairwise by the Diebold Mariano test. Here the quadratic error are used in the test which puts more emphasis on big errors in prediction. The absolute value or even non-symmetric transformations of errors are other options. Table 4 gives the results of the tests. Bold font indicates a significant difference at the 5% level and a positive number indicates the column model outperforms the row model.

To explore how the models are chosen across time, Figure 3 graphs model complexity based on particular criteria for each model. For the Lasso and Elastic net choosing more parameters to be non-zero amounts

Table 4: Comparison of Monthly Out-of-Sample Prediction using Diebold-Mariano Tests

	Ridge	Lasso	ENET	PCR	PLSR	GLM	RandomF	GBRT
OLS	2.40	2.03	1.95	1.64	0.63	0.10	1.95	1.93
Ridge		-1.57	-1.84	-1.97	-1.78	-1.62	-2.59	-0.89
Lasso			-2.18	-2.11	-1.46	-1.27	-1.59	0.28
ENET				-1.79	-1.41	-1.22	-1.44	0.51
PCR					-0.80	-0.85	-0.38	1.44
PLSR						-0.38	0.53	1.43
GLM							0.85	1.23
RandomF								1.33

Note: This table reports pairwise Diebold-Mariano test statistics comparing the out-of-sample prediction performance among nine models. Positive numbers indicate the column model outperforms the row model. Bold font indicates the difference is significant at 5% level or better.

to more complexity. For PCR and PLSR the no. of directions is chosen as a measure. For the GBRT the chosen tree depth is plotted. Finally for the Random Forest the minimum terminal node size is plotted. Note that the Random Forest is inverted in terms of complexity as larger node size means a less complex tree.

Note: This figure demonstrates the model complexity over time for Lasso, Elastic net (ENet), PCR, PLS, Random Forest (RF) and Gradient Boosted Regression Trees (GBRT) in each training sample of the 18-year recursive out-of-sample analysis. For ENet and Lasso the number of variables selected to have non-zero coefficients are reported; for PCR and PLS reported is the number of selected components/directions; for RF reported is the tuned minimal node size; and for GBRT the tuned tree depth.

There is a bit of a tendency for the model complexity to decrease over time. An explanation for this could be that the “information” available in the marked have decreases as more of the computational techniques have been adopted over time.

5 Conclusion

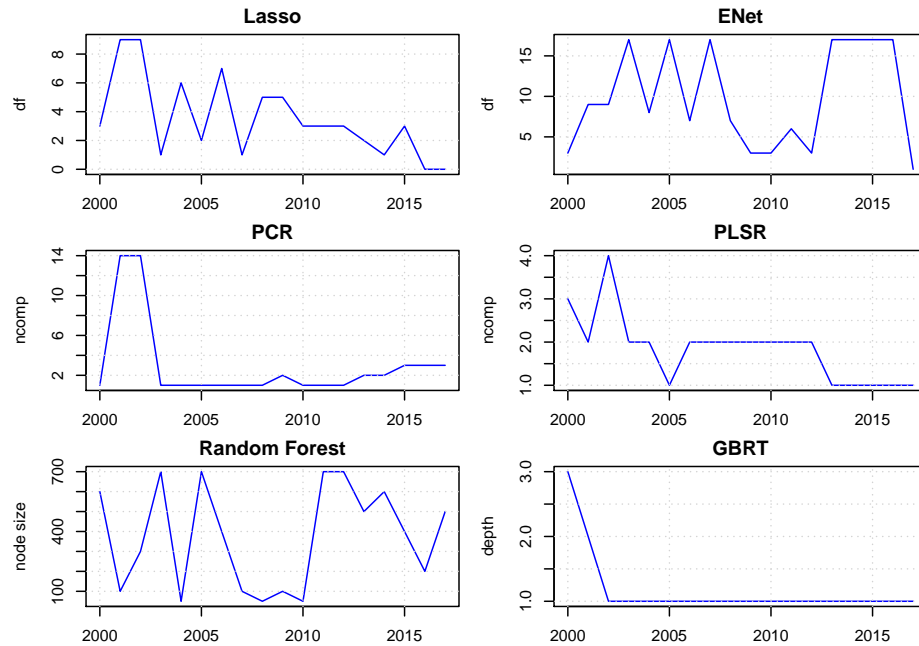
The main finding in this paper is that Machine Learning methods can pose an important tool in empirical asset pricing and financial/economic forecasting. Both simpler methods like Lasso regression and more advanced techniques like Random Forests provide more stable methods for forecasting, even in high noise data. The Lasso improves on the linear model by shrinking and selecting parameters in a continuous way by adding a penalty to the objective function. This improves on the models variability, even on the linear model with variable selection (see Friedman, Hastie, and Tibshirani 2001, vol. 1, chap. 3.4). Random Forests presents a complete alternative method to the linear model by using tree regression for the conditional expectation. Together with the gradient boosted tree regression, these methods are the standout performers in this high noise environment, especially with the non-linear data generating process.

There is an ambiguity however Something about ridge performing well in empirical exercise...

6 References

- Cybenko, George. 1989. “Approximations by Superpositions of a Sigmoidal Function.” *Mathematics of Control, Signals and Systems* 2: 183–92.
- Diebold, Francis X, and Robert S Mariano. 2002. “Comparing Predictive Accuracy.” *Journal of Business & Economic Statistics* 20 (1). Taylor & Francis: 134–44.

Figure 3: Model complexity over time



Friedman, Jerome, Trevor Hastie, and Robert Tibshirani. 2001. *The Elements of Statistical Learning*. Vol. 1. 10. Springer series in statistics New York.

Gu, Shihao, Bryan Kelly, and Dacheng Xiu. 2018. “Empirical Asset Pricing via Machine Learning.” Working Paper 25398. Working Paper Series. National Bureau of Economic Research. <https://doi.org/10.3386/w25398>.

Hamilton, James D. 1995. *Time Series Analysis. Economic Theory. II*, Princeton University Press, USA.

James, Gareth, Daniela Witten, Trevor Hastie, and Robert Tibshirani. 2013. *An Introduction to Statistical Learning*. Vol. 112. Springer.

Kelly, Bryan T, Tobias J Moskowitz, and Seth Pruitt. 2018. “Understanding Momentum and Reversal.” Available at SSRN 3269897.

Masters, Timothy. 1993. *Practical Neural Network Recipes in C++*. Morgan Kaufmann.

Newey, Whitney K., and Kenneth D. West. 1987. “A Simple, Positive Semi-Definite, Heteroskedasticity and Autocorrelation Consistent Covariance Matrix.” *Econometrica* 55 (3). [Wiley, Econometric Society]: 703–8. <http://www.jstor.org/stable/1913610>.

Stocks, Bonds. 1995. “Bills and Inflation 1995 Yearbook.” Ibbotson Associates, Chicago.

Welch, Ivo, and Amit Goyal. 2008. “A Comprehensive Look at the Empirical Performance of Equity Premium Prediction.” *The Review of Financial Studies* 21 (4). [Oxford University Press, Society for Financial Studies]: 1455–1508. <http://www.jstor.org/stable/40056859>.

Zou, Hui, and Trevor Hastie. 2005. “Regularization and Variable Selection via the Elastic Net.” *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 67 (2). Wiley Online Library: 301–20.

7 Appendix

Table 5: Comparison of Predictive R^2 s for Machine Learning Algorithms in Simulations (unaltered)

DGP	(a)				(b)			
Parameter	$P_c = 50$		$P_c = 100$		$P_c = 50$		$P_c = 100$	
$R^2(\%)$	IS	OOS	IS	OOS	IS	OOS	IS	OOS
OLS	10.11	-1.30	12.90	-3.79	6.87	-5.73	9.45	-8.82
Ridge	8.06	1.47	9.35	1.01	4.56	-0.19	5.54	-0.15
Lasso	8.12	2.61	9.37	1.86	4.54	-0.43	5.60	-0.03
ENet	8.18	2.69	9.53	1.89	4.66	-0.41	5.72	-0.07
PCR	7.46	0.08	8.32	-0.04	4.23	-0.68	4.97	-0.72
PLSR	7.38	0.27	8.51	-0.16	3.67	-1.98	4.50	-3.36
GLM	7.92	0.21	9.07	0.01	4.80	-0.48	5.87	-0.98
RandomF	13.91	2.33	16.39	1.83	13.49	1.02	14.74	2.23
GBRT	9.84	1.94	11.59	0.55	10.06	1.09	10.37	1.20
NN1	-640.58	-1217.57	-74.98	-40.74	-198.00	-411.69	-38.00	-11.00
NN2	-20676.24	-6461.09	-7219.14	-2839.42	-11558.00	-3138.64	-16728.12	-4427.85
NN3	-28342.48	-5195.43	-26138.54	-4083.54	-30438.74	-5148.84	-19274.67	-4973.21
Oracle	8.41	3.45	10.04	3.74	4.89	0.12	5.81	-0.07

Note: Table 1 except all Monte Carlo repetitions are kept.